

Stealthy and Efficient Adversarial Attacks against Deep Reinforcement Learning (SEAADRL) – AAAI 2020

Jianwen Sun Tianwei Zhang
Xiaofei Xie Lei Ma Yan Zheng Kangjie Chen Yang Liu

Presented by Aayush Naik

Table of Contents

- 1 Motivation
- 2 Background
- 3 SEAADRL
- 4 Experiments and Results
- 5 Extension

Table of Contents

- 1 Motivation
- 2 Background
- 3 SEAADRL
- 4 Experiments and Results
- 5 Extension

Motivation

Analyzing attacks against DL & DRL important because:

- DL & DRL are becoming widely used.
- Especially in high-stakes and safety-critical scenarios:
autonomous driving,
weapon systems,
medical assistance.



Table of Contents

- 1 Motivation
- 2 Background
- 3 SEAADRL
- 4 Experiments and Results
- 5 Extension

The RL Setting

The environment is modeled as a Markov Decision Process (MDP) [Bel57], given by the tuple $(\mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma)$.

- \mathcal{S} : a discrete/continuous space of states.
- \mathcal{A} : a discrete/continuous action space.
- $\mathcal{P} : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow [0, 1]$: transition function.
- $\mathcal{R} : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$: reward function.
- γ : discount factor.




Figure: The RL Setting [SB18]

Attacks Against Supervised DL and DRL

Since the work of Szegedy et. al [SZS⁺13], several attacks have been demonstrated on supervised deep learning.





Figure: Ostrich (left), also ostrich (center), ostrich (right). Image credits: [SZS⁺13].

Attacks Against Supervised DL and DRL

Two types of attacks:

- Untargeted attacks: misclassification to an arbitrary class/label [SZS⁺13, MDFF16].
- Targeted attacks: misclassification to a *specific* class/label [CW17].

Moving to DRL (initial works):

- Apply attacks from supervised DL to inputs of DRL agents [HPG⁺17].

Key Ideas

- 1 Targeted attacks against supervised DL can be used to make DRL agents choose specific suboptimal actions [HGP19].
- 2 Going for a small number of *critical hits*, instead of attacking at every step. *Timing* attacks to deal the maximum “damage” in the least number of steps [KS17, SZX⁺20].

Table of Contents

- 1 Motivation
- 2 Background
- 3 SEAADRL**
- 4 Experiments and Results
- 5 Extension

Two components of a targeted DRL attack:

- 1 *How*: Generating noise added to the input (perturbation) to make the agent choose an action.
- 2 *When*: Deciding when to do perturbation to maximize “damage” in the least number of steps.

SEAADRL uses existing methods for the *how* [CW17], and its main contribution is in the *when*. Two algorithms/classes of attack:

- 1 Critical-point (CP) attack.
- 2 Antagonist attack.

SEAADRL: Critical-point Attack

- At every timestep t , consider all possible N -step action sequences.
- Compute the *divergence* of each of “final” states.
- Pretend that we have access to a transition function of the environment: $\hat{\mathcal{P}}$.




Figure: $\mathcal{A} = \{a, b\}$ and $N = 2$ in this example. $T(s_{t+2})$ is the divergence of state s_{t+2} .

SEAADRL: Critical-point Attack: Divergence Function

- User-defined domain specific function.
- For TORCS,

$T(s)$:= distance between car and center of road.




Figure: Low divergence; near center.

Figure: High divergence; off center.

SEAADRL: Critical-point Attack

- Let our trained DRL agent have policy π .
- Compare the divergence of the N -step action sequences with the *baseline*.
- Perform attack if $|T(s_{t+2}) - T(s'_{t+2})| > \Delta$. Here, $N = 2$.

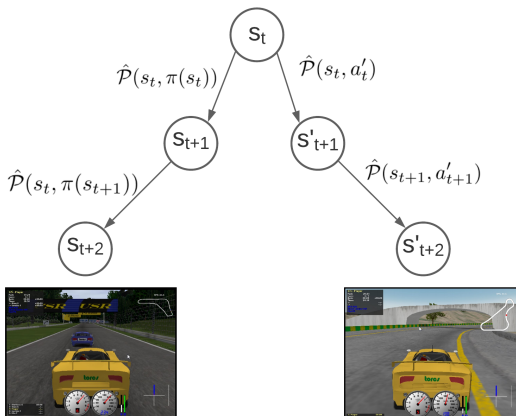


Figure: Baseline expected state (left); highest-divergence action sequence (right).

SEADRL: Critical-point Attack: Summary

- Train environment transition model $\hat{\mathcal{P}}$.
- At every timestep t , look ahead N steps with all possible action sequences.
- If any of those sequences lead to high divergence, perform that attack.
- Attack only once per episode (efficiency).
- Big disadvantage: highly domain specific.

SEAADRL: Antagonist Attack

- Train an agent (called antagonist) that receives $-r_t$ at every time step.
- Allow this agent to decide whether to attack or not at each timestep (when).
 - Limit number of antagonist actions (attacks) to a small number of steps (2-5) per episode.
 - A regular, pretrained DRL agent acts for other steps.
- The antagonist learns to pick opportune moments to act and do significant damage.
- Main advantage over CP attack: requires no domain knowledge.

Table of Contents

- 1 Motivation
- 2 Background
- 3 SEAADRL
- 4 Experiments and Results**
- 5 Extension

Experiments and Results

- The authors of SEAADRL run their experiments:
 - with different DRL algorithms (A2C/A3C, PPO and DDPG).
 - on many environments (Atari Pong, Atari Breakout, TORCS, MuJoCo HalfCheetah and more).
- Show significant decline in average returns with attacking only a very small number of steps per episode.
- My replication only on Atari Pong with A2C.

My Replication: Plots





Figure: Average return vs. Δ for CP attack





Figure: Average return vs. # attack steps for antagonist attack

My Replication: Summary

- Smaller scale of experiments.
- Overall agreement with original experiments.
- Minor disagreements with details.

Table of Contents

- 1 Motivation
- 2 Background
- 3 SEAADRL
- 4 Experiments and Results
- 5 Extension**

Extension: Robust Training with Antagonist Agent

Main Idea: While training a regular DRL agent (e.g. A2C), use a trained antagonist agent to inject damaging actions at critical moments to make training more robust.

Agent/Algorithm	Mean Return	Return Std-Dev
A2C	20.23	0.85270
A2C with Robust Training	21	0.0

Table 1: Results of preliminary experiments for Robust agent.






Intuition on why this works: our new agent will avoid states that can lead to significant damage within a small number of steps, and thus stay in a “safer space”.

Further: Do this iteratively.


Thank You!

Thank You!

References I

-  Richard Bellman, *A markovian decision process*, Journal of mathematics and mechanics (1957), 679–684.
-  Nicholas Carlini and David Wagner, *Towards evaluating the robustness of neural networks*, 2017 iee symposium on security and privacy (sp), IEEE, 2017, pp. 39–57.
-  Léonard Hussenot, Matthieu Geist, and Olivier Pietquin, *Targeted attacks on deep reinforcement learning agents through adversarial observations*, arXiv preprint arXiv:1905.12282 (2019).
-  Sandy Huang, Nicolas Papernot, Ian Goodfellow, Yan Duan, and Pieter Abbeel, *Adversarial attacks on neural network policies*, arXiv preprint arXiv:1702.02284 (2017).
-  Jernej Kos and Dawn Song, *Delving into adversarial attacks on deep policies*, arXiv preprint arXiv:1705.06452 (2017).

References II

-  Seyed-Mohsen Moosavi-Dezfooli, Alhussein Fawzi, and Pascal Frossard, *Deepfool: a simple and accurate method to fool deep neural networks*, Proceedings of the IEEE conference on computer vision and pattern recognition, 2016, pp. 2574–2582.
-  Richard S Sutton and Andrew G Barto, *Reinforcement learning: An introduction*, MIT press, 2018.
-  Christian Szegedy, Wojciech Zaremba, Ilya Sutskever, Joan Bruna, Dumitru Erhan, Ian Goodfellow, and Rob Fergus, *Intriguing properties of neural networks*, arXiv preprint arXiv:1312.6199 (2013).



Jianwen Sun, Tianwei Zhang, Xiaofei Xie, Lei Ma, Yan Zheng, Kangjie Chen, and Yang Liu, *Stealthy and efficient adversarial attacks against deep reinforcement learning*, Proceedings of the AAAI Conference on Artificial Intelligence, vol. 34, 2020, pp. 5883–5891.